


Utility versus measured computation: six matched learned policies



Same training episodes, five fits and 30 paired evaluation seeds per scenario. One workstation; engine and I/O excluded.

MAP uses a direct dot product; Bayes includes quadrature. This is not a browser-model or end-to-end benchmark.