







Separating stay/change mass from alternative scores

Held-out services; saved distributions only. All three seeds. Original training study: FAIL (18/22).



Held-out-service accuracy (all 7,819 rows; seed means): flat 93.36% | MM 92.45% | AA 91.81% | MA 92.51% | AM 91.62%

MM/AA reconstruct the original models. MA/AM combine two source distributions; both source predictions have a cost.

Exposed TRAIN development with correct previous values supplied. Seeds repeat optimization, not independent service sampling.

Probability-factor diagnostic, not a trained model, causal gate intervention, fresh evaluation, or reversal of the failed study.