


Qwen observation baseline with a correct previous value

Exposed official TRAIN | 578 changed and 7,241 retained rows per arm | Zoomed metric axes

Current vs historical flat: behavior FAIL; proper scores FAIL

History vs current: behavior FAIL; proper scores FAIL



Historical points are three earlier corrected-flat fits, not repeated Qwen runs. No probability-calibration or architecture claim.