


Token alignment: behavioral continuation FAIL

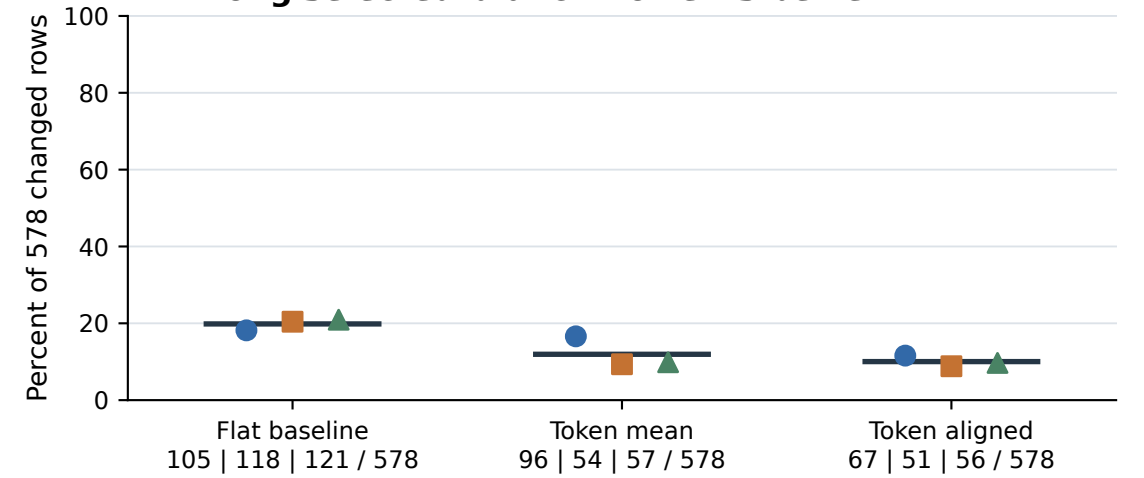
Historically exposed TRAIN development split; correct previous state supplied. All nine final fits and both controls shown.

6201 6202 6203 Three-seed mean in A/B Three-seed mean in C/D

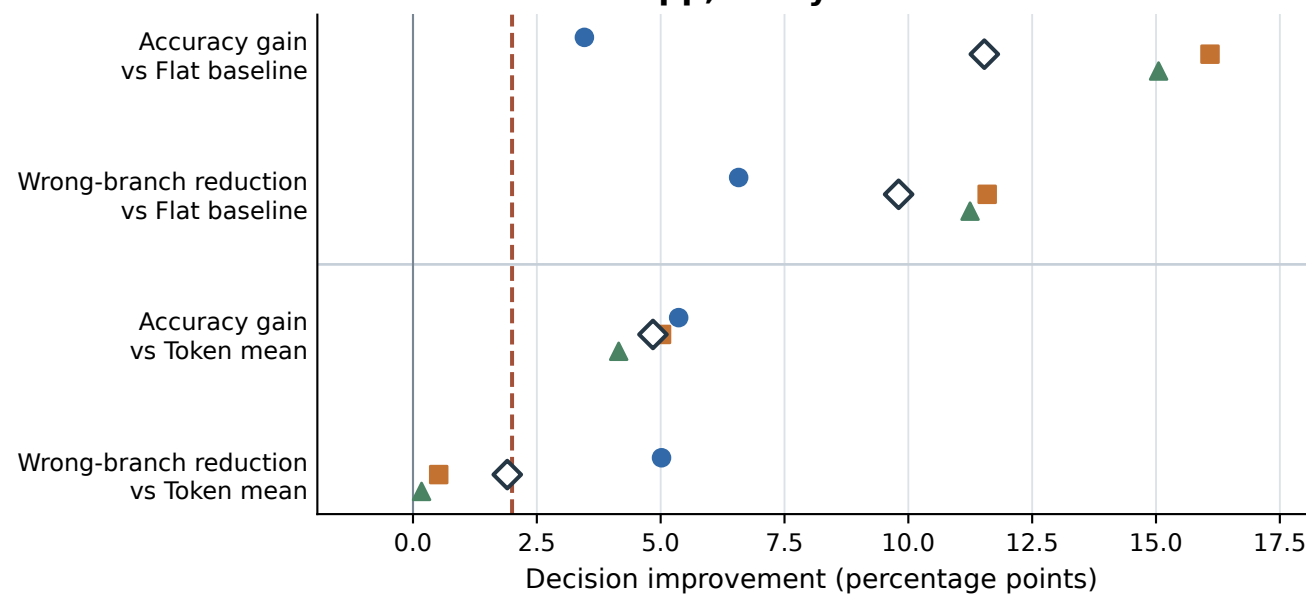
A Changed-state accuracy: higher is better



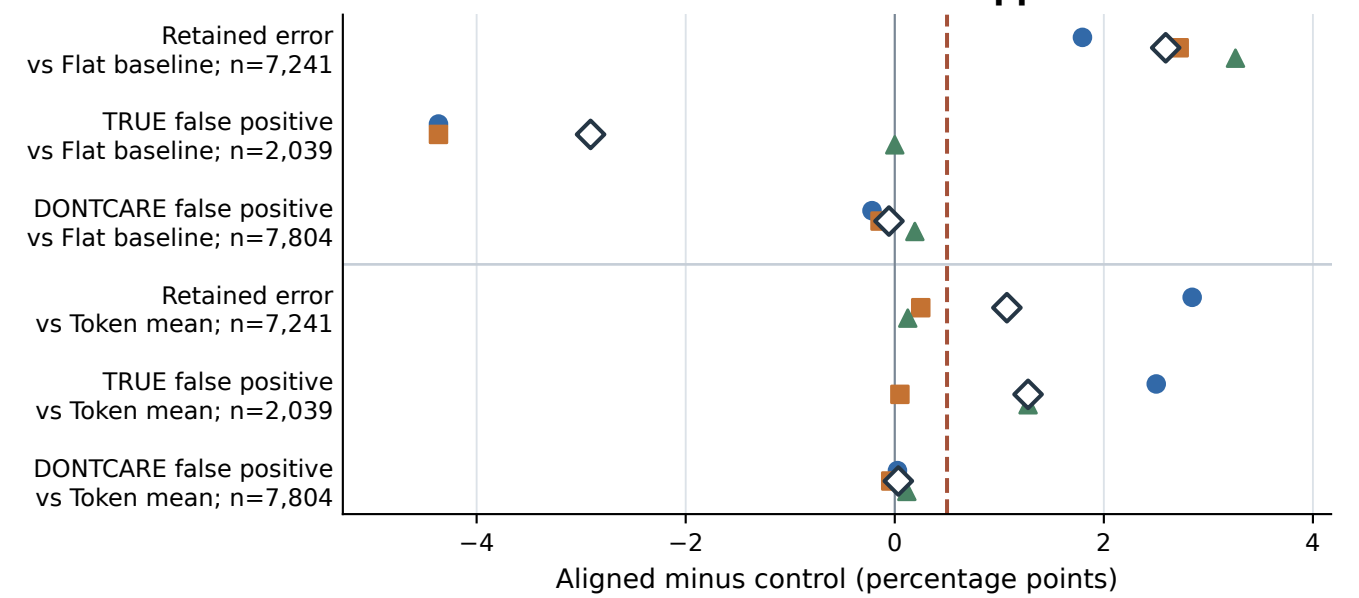
B Wrong selected branch: lower is better



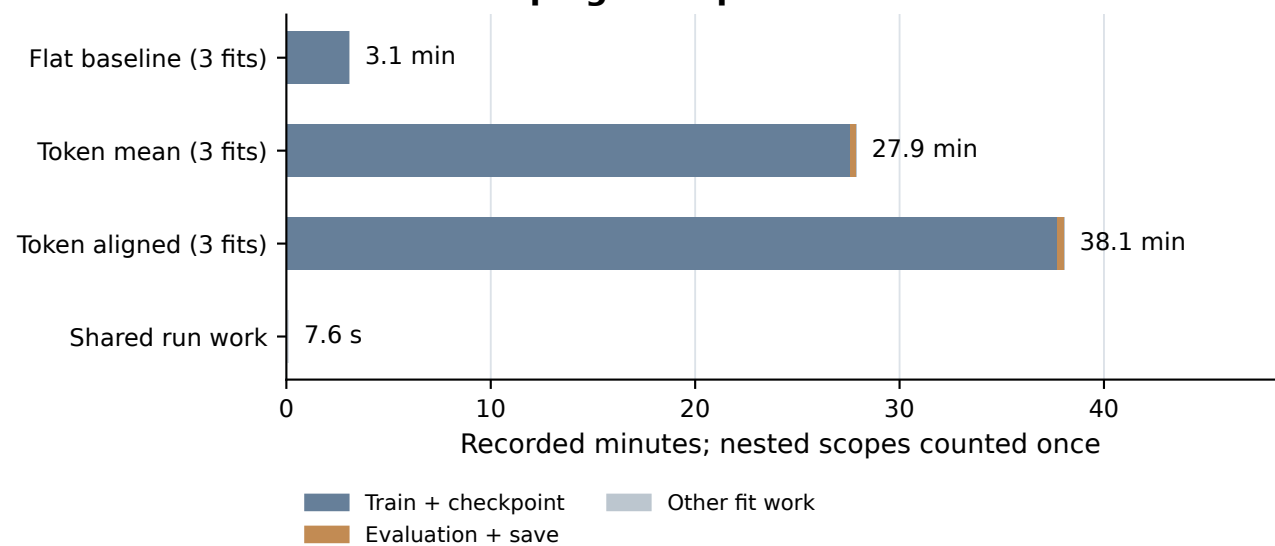
C Gains: mean ≥ 2 pp; every seed ≥ 0



D Harm limits: mean increase ≤ 0.5 pp



E Actual campaign compute



F Scope and fixed rule

Both controls required: 18/22 checks passed
Against flat baseline: 10/11
Against token mean: 8/11

Primary panel: 6 held-out services, 7,819 rows
Changed: 578 | Retained: 7,241
TRUE changes: 29 | DONTCARE changes: 5
No FALSE changes or clears; sparse recalls descriptive.

Whole run: 69.19 min / 120 min cap
Peak RSS: 1.74 GiB (process lifetime)
Schema preparation: 4.99 s, separate

A/B count labels follow seed order 6201 | 6202 | 6203. Branch is the branch of the selected candidate, not the largest summed branch mass.

Three seeds repeat optimization, not independent service sampling. No novelty, calibration, deployment or fresh-confirmation claim.

The earlier 48-minute capacity admission failed and remains closed; this campaign used a separately declared 120-minute allocation.