


Typed output and rare-category support: FAIL (5/9 fixed checks)

Historically exposed internal TRAIN service split; correct previous value supplied. All 12 final fits shown.

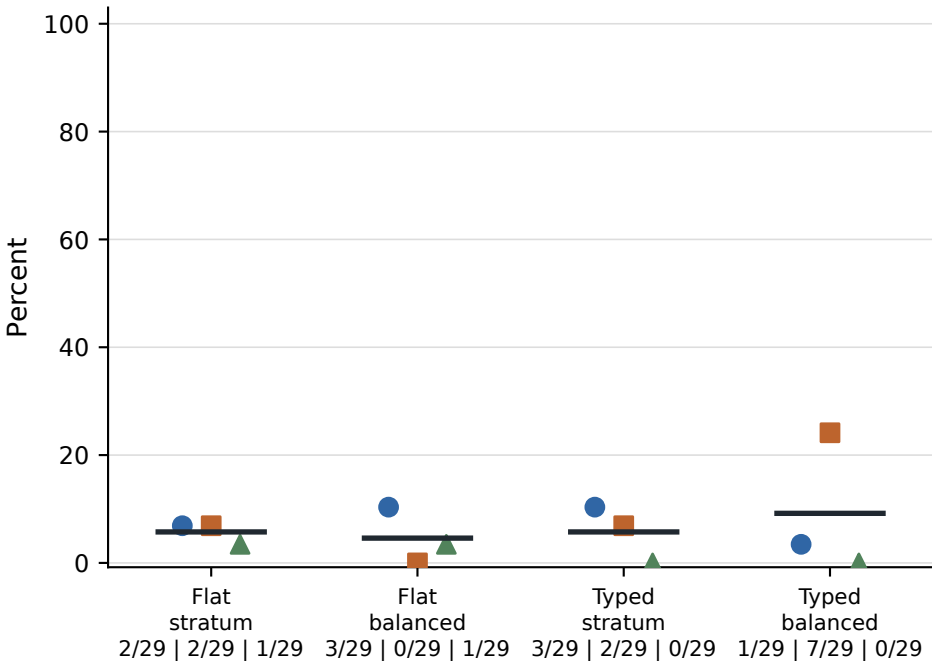
6101 6102 6103 Three-seed mean

A Primary paired comparison



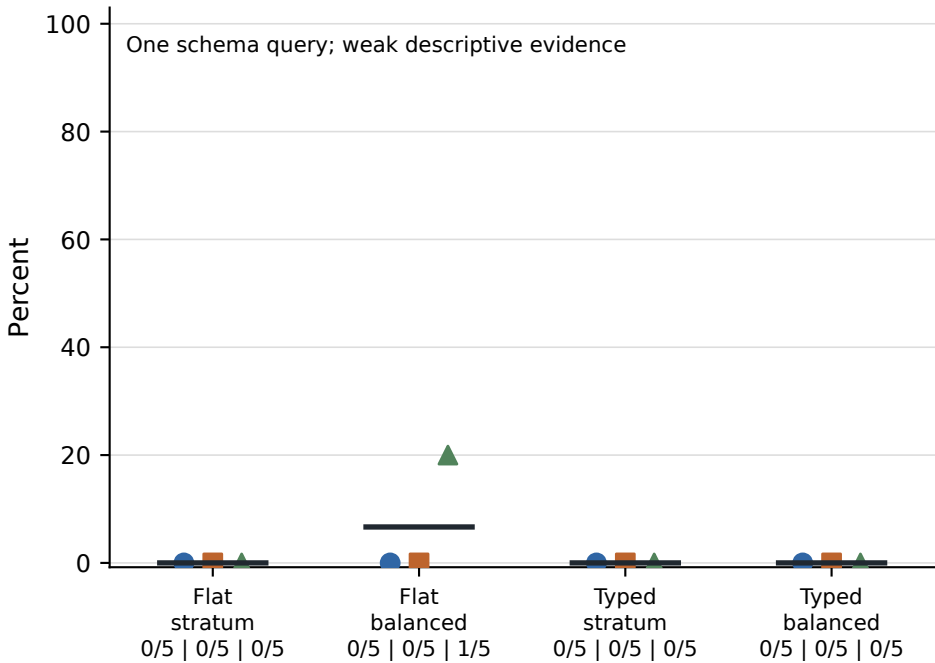
578 changed rows; lower is better.
Every paired seed must be nonworse.

B TRUE-change recall: 29 rows



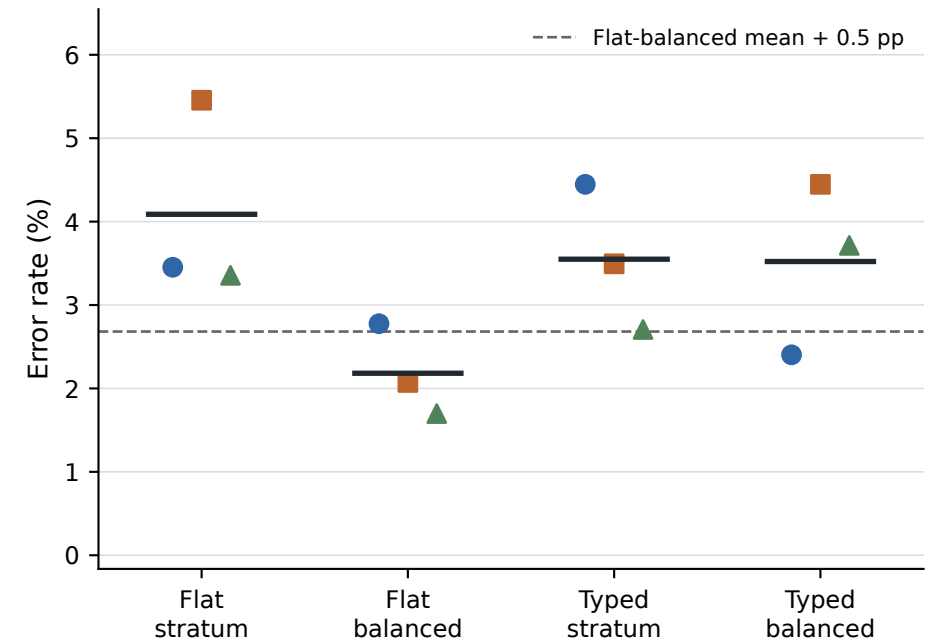
Correct/total counts follow seed order 6101 | 6102 | 6103.

C DONTCARE recall: only 5 rows



Correct/total counts follow seed order 6101 | 6102 | 6103.

D Retained-state error

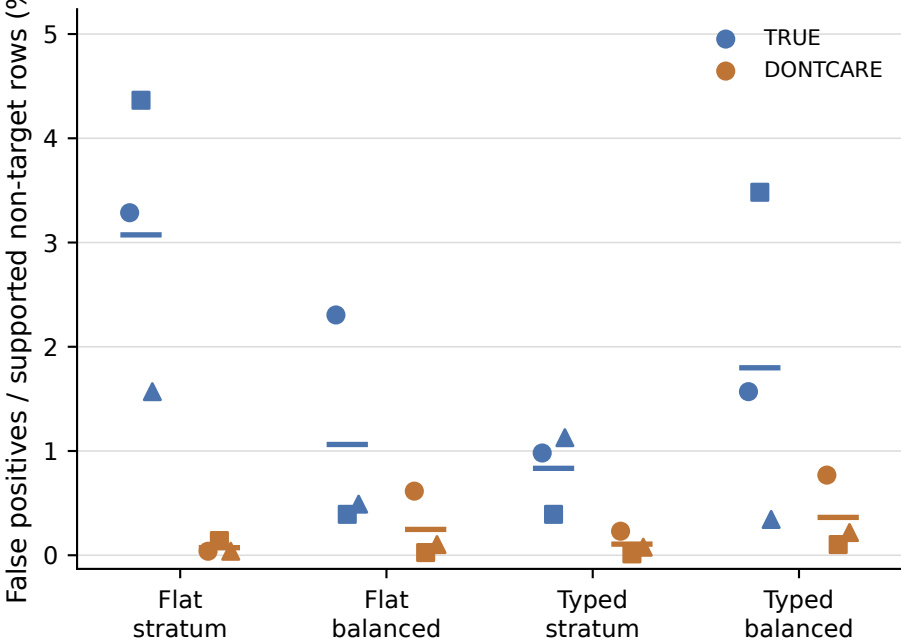


Whole run: 551.8 s; process-lifetime peak RSS: 1.79 GiB. Fit time is not inference latency.

Three seeds are repeated optimizations, not independent service samples. No FALSE/clear changes in the primary panel.

Inherited feature preparation and separate preflight/reporting are additional costs. This figure does not establish calibration, recurrence, or fresh confirmation.

E Rare-type false positives



Lower is better; allowed mean increase is 0.5 pp per type.
Denominators exclude rows where that candidate type is absent.

F Actual paid fit cost

