





Can direct action heads preserve search utility and reduce complete computation?

1,536 autonomous episodes | All 12 fitted heads and 4 unchanged planners | 48 paired cases per regime

Full-belief head competence: FAIL Pilot: FAIL Candidate checks: 6/40 passed No architecture novelty established

Solid rows: fitted seeds 7901/7902/7903. Hatched rows: original analytic planner. Every fitted seed remains visible.



Both regimes use their own initial-hit mixture. Found counts are unweighted. All failed searches contribute the full 2,188-move horizon.

Controller time includes feature construction, standardization, inference/planning, legal selection, updates and allocated setup; one rotated CPU pass.

Training and collection cost are separate. Evolving actor state excludes immutable learned weights/standardizers, shared tables and temporary workspaces.

Full belief has 362,788 head parameters versus 199,396 for compact/recent heads. Matching hidden widths does not equalize model capacity.

Known lambda 4 and longer horizons are transfer tests. No seed/fill selection, prior-gate reversal or learned-architecture claim.