

OpenJev: Bounded Decisions, Causal Memory, and Doom Control

OpenJev project

<https://github.com/kw2828/OpenJev>

September 16, 2026

Working research draft; not a conference submission

Abstract

OpenJev exposes request-defined candidate decisions and a reproducible Doom control testbed. An initial 440-episode Bayesian pilot did not establish a control gain from posterior averaging. We then froze a matched causal-history ablation using the same 80 historical training episodes, five fitted replicates, and 1,980 new evaluation and audit episodes. A 17-feature history head improved mean net utility over the original seven-feature MAP head by 1.779 in Defend the Center and 7.301 in Defend the Line; multiplicity-adjusted bootstrap intervals had positive lower endpoints. However, both comparisons against a 17-feature current-only control remained inconclusive. The predeclared continuation gate therefore failed, and adaptive-compute and second-environment efficacy experiments were not launched. Separate original gameplay and score-kernel measurements are retained with their own protocols. We release causal feature code, frozen protocols, traces, figures, and a browser-local decision playground. The results establish a local improvement over the original model, but do not isolate a memory-specific advantage over richer current-state features, demonstrate a new RL algorithm, or establish superiority over rules.

1 Motivation and scope

A bounded application decision often needs a stable identifier and a score rather than generated prose. OpenJev accepts text context, a set of questions, and candidate identifiers with English descriptions. The application maps the returned identifier to an action. A pretrained Qwen language backend scores next-token candidate labels; it remains an autoregressive transformer and is not a new architecture. Questions are processed sequentially, without shared-prefix optimization.

The project is inspired by the typed-decision interface of Jev. It does not reproduce Jev’s weights or proprietary Reinforcement Learning for Calibrated Decisions (RLCD). The research model in this report is a separate seven-feature Bayesian logistic outcome predictor. There is no language-model RL training, PPO, GRPO, or Bellman update in the pilot.

The initial pilot asks: *does approximate Bayesian averaging or a conservative posterior bound improve firing decisions relative to the point estimate of the same learned model?* The original language-controller recording is a demonstration, not one of the fitted models or a baseline in this experiment.

2 Decision interface and outcome model

For candidate logits $z_1, \dots, z_K$, the interface returns the conditional choice distribution

$$q(a_i \mid c, Q, C) = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)}. \quad (1)$$

Here c is context, Q is the question, and C is its candidate set. Candidate identifiers are preserved by the runtime. These uncalibrated scores sum to one but are not probabilities of task success. The runtime also reports candidate-label mass under the full vocabulary; restricting the denominator to candidates does not eliminate omitted-answer uncertainty.

The outcome model instead predicts a specified binary event when a fire command is issued. Let $x \in \mathbb{R}^7$ contain an intercept, visibility, clipped absolute aim error, target half-width, clipped distance, scaled health, and clipped ammo. With $y = 1$ for a hit-count increase during the following seven game tics,

$$p(y = 1 \mid x, w) = \sigma(x^\top w), \quad w \sim \mathcal{N}(0, I), \quad (2)$$

$$\hat{w} = \arg \min_w \sum_n \left[\log(1 + e^{x_n^\top w}) - y_n x_n^\top w \right] + \frac{1}{2} \|w\|^2, \quad (3)$$

$$\Sigma = \left[I + X^\top \text{diag}(p_n(1 - p_n))X \right]^{-1}. \quad (4)$$

Newton steps with line search fit the MAP estimate, with at most 40 iterations. A Laplace approximation uses $w \mid D \approx \mathcal{N}(\hat{w}, \Sigma)$. Twenty-node Gauss–Hermite quadrature approximates $\bar{p}(x) = \mathbb{E}_{w \mid D}[\sigma(x^\top w)]$. The lower probability used for conservative gating is

$$p_{\text{low}}(x) = \sigma \left(x^\top \hat{w} - 1.645 \sqrt{x^\top \Sigma x} \right). \quad (5)$$

This is approximately a one-sided 95% posterior lower bound conditional on the fitted approximation, not a conformal or frequentist coverage guarantee. The implementation separates $\text{Var}_{w \mid D}(p)$ and $\mathbb{E}_{w \mid D}[p(1 - p)]$. Neither decomposition corrects feature misspecification or temporal dependence.

The three learned policies issue fire when p_{MAP} , $\bar{p}$, or p_{low} respectively exceeds the fixed threshold 0.25. This is a myopic decision rule; it does not account for future survival or information value. A negative Brier reward primitive is available in the package, but is *not* the fitting loss used here.

3 Experimental design and measurement

We use ViZDoom (Kempka et al., 2016), with privileged visible-actor labels and game variables rather than pixels. Five independent collection seeds each produce 16 training episodes in Defend the Center. The behavior policy requests fire with probability 0.6 and randomizes steering with probability 0.2. Only issued fire-command windows provide outcome labels; cooldown windows without a hit count as failures. Waiting does not generate a synthetic negative label.

All five evaluated firing policies share rule-based steering. The comparisons are the existing rule, the tiny imitation model’s firing output, logistic MAP, Bayesian mean, and Bayesian lower bound. Each learned policy is evaluated on ten paired seeds per scenario for each of five training replicates. The rule and imitation baselines are each run once per seed and reused in paired comparisons, not counted as independent copies. Evaluation seeds 40000–40009 are disjoint from fitting. Episodes have at most 180 seven-tic decisions, and there is no evaluation-time fitting or threshold tuning.

The 440 episodes comprise 80 training episodes, 20 common-audit episodes and 340 evaluation episodes. The common audit applies a randomized behavior policy to supply the same observed firing events to all fitted models. It has ten episodes per scenario, containing 349 center and 700 line events. These events are dependent within episodes and do not expose every counterfactual action outcome. Utility is

$$U = \sum_{t: \text{issued fire}} (y_t - 0.25). \quad (6)$$

It measures command-window outcomes, not per-bullet accuracy or the default game reward. Intervals use 10,000 paired crossed-bootstrap resamples over five training replicates and ten evaluation seeds. They are exploratory, unadjusted intervals with small sampling units.

Preserved failed measurement. The preceding v1 run also completed 440 episodes, but used net ammo decrease to detect shots. Ammo replenishment in Defend the Line masked firing, causing zero detected shots and empty calibration audits. The original traces and quality exception are retained. V2 changes the estimand to commanded firing windows and uses fresh evaluation seeds; it is a new development study, not a repaired identical replication or sealed confirmation. Only v2 supports the cross-scenario tables here.

4 Bayesian pilot results

Controller	Center utility	Line utility	Line kills
Rule baseline	3.175	13.900	26.90
Tiny imitation	3.150	13.100	25.40
Logistic MAP	0.855	5.715	16.34
Bayesian mean	0.870	5.140	16.96
Bayesian lower bound	1.000	0.575	4.76

Table 1: Mean v2 outcomes. Learned controllers have 50 evaluation episodes per scenario; each existing baseline has ten. All policies share rule steering. The table is generated directly from the preserved analysis JSON.

Bayesian mean minus MAP utility was $+0.015$ in the center scenario (exploratory 95% interval $[-0.705, 0.725]$), and -0.575 in the line scenario ($[-3.745, 1.960]$). These intervals do not establish a benefit, equivalence, or absence of smaller effects. The conservative lower bound performed worse in the line scenario: -5.140 relative to MAP ($[-9.120, -1.720]$). Strong rule and imitation baselines further limit any claim of learned-control progress.




Figure 1: Paired utility differences from MAP. Positive values favor the compared controller. Intervals account for training replicate and episode-seed variation, but are exploratory and unadjusted.




Figure 2: Paired audit Brier scores across five fitted models. The same events are reused across fits, so the five lines are not five independent audit datasets. Red diamonds are means, not confidence intervals.

On the shared line audit, the mean Brier score decreased descriptively from 0.230733 to 0.226479 with Bayesian averaging; in the center it changed from 0.122786 to 0.122836. Brier score measures overall probability accuracy and is not an isolated calibration diagnostic. No uncertainty interval is reported for this descriptive audit comparison. Its apparent improvement in the shifted scenario did not coincide with better mean control utility. This local result motivates measuring prediction and decision performance separately, rather than treating one as a substitute for the other.

5 Frozen causal-memory follow-up

The initial head omitted temporal state. We therefore tested a history extension against both the original seven-feature model and a 17-feature current-only expansion. The protocol and runner were committed at 04ce32d before evaluation. All variants reuse exactly the same 80 historical training episodes and commanded-fire labels. The five original MAP weight vectors were reproduced exactly, ruling out changed fitting data as the source of the comparison.

History appends the previous six non-intercept observation features, previous issued-fire and observed-hit indicators, time since the last issued fire command (capped at ten windows and scaled), and a has-previous indicator. History resets per episode, is updated only after action feedback, and contains no current-window outcome or hidden weapon readiness. The current-only expansion instead appends squares and four interactions. It matches feature count, but does not ensure identical effective capacity or regularization geometry. Both MAP and posterior mean are evaluated for each representation.

The new run uses untouched evaluation seeds 50000–50029 and separate audit seeds 60000–60029 in each scenario. It comprises 1,920 policy evaluations and 60 common-audit episodes; reused training episodes are not counted as new executions. All policies share rule steering, seven-tic windows and a cost of 0.25 per issued fire window. The episode cap remains 180 decisions. Some line episodes hit this cap, so results concern a capped horizon. Evaluation job order is randomized.

Continuation rule. For each scenario, history MAP must improve over both original MAP and expanded MAP by at least 0.5 mean net utility, with a positive lower endpoint of a Bonferroni-adjusted

98.75% interval. Net utility is raw utility minus one unit per second of accumulated decision computation. Four primary contrasts use 20,000 paired crossed-bootstrap draws over five fits and 30 seeds. These are finite-sample exploratory intervals, not a guaranteed test or external confirmation. History must also use no more than three times the original model’s aggregate decision time, with every episode’s p95 decision time below 1 ms.

Timing includes feature construction, steering, prediction, decision construction and history updates. It excludes engine observation/advance and trace I/O. MAP uses a direct dot product; posterior averaging includes quadrature. Lightweight document editing occurred concurrently on the workstation; browser GPU inference began only after this run ended. The timing screen is local, not a controlled systems benchmark.

Policy	Center U	Line U	Center ms	Line ms
Current MAP	0.823	6.493	0.010	0.010
Expanded MAP	1.012	4.845	0.019	0.018
History MAP	2.603	13.795	0.020	0.019
Current Bayes	0.683	5.213	0.023	0.021
Expanded Bayes	1.027	6.205	0.030	0.028
History Bayes	2.837	12.680	0.030	0.030
Rules	2.875	15.233	0.005	0.005
Imitation	2.917	14.817	0.028	0.024

Table 2: New memory-ablation results: raw command-window utility U and mean decision time in milliseconds. Each learned policy has 150 episodes per scenario; each fixed baseline has 30. Rules and imitation remain strong controls.

History MAP improved net utility over original MAP by +1.779 in Center (adjusted interval [0.277, 3.564]) and +7.301 in Line ([0.351, 14.180]). Against expanded MAP, the differences were +1.591 ([−0.349, 3.571]) and +8.950 ([−0.972, 16.998]). Both timing gates passed: history used approximately 1.92 and 1.94 times the original per-decision compute, and maximum episode p95 times were 0.135 and 0.123 ms. The two capacity-control intervals nevertheless include zero. Thus the full continuation gate failed.




Figure 3: All four predeclared memory contrasts. A positive mean alone is insufficient for continuation; both capacity-control intervals cross zero. The dotted vertical line marks the minimum required mean improvement of 0.5.




Figure 4: Measured utility and decision compute for the six learned policies. These descriptive points do not establish a statistically dominant frontier. Gate decisions use the paired primary contrasts and timing limits, not visual selection of the best point.

Shared-audit mean MAP Brier scores changed from 0.131446 to 0.119535 in Center and 0.243435 to 0.228333 in Line with history. These descriptive probability-accuracy results do not replace the primary control comparisons. No improvement uniquely attributable to Bayesian inference or learned recurrent computation is established. Adaptive-compute and second-environment efficacy runs were not launched, as required by the frozen failure rule; no new tuning run was used to obtain a pass. The evidence suggests weakness in the original memoryless specification but does not isolate history as the cause relative to the richer current-only representation.

6 Separate gameplay and systems benchmarks

Original imitation benchmark. The shipped 5,253-parameter MLP uses 11 numeric features and two 64-unit hidden layers, with separate steering and firing heads. It was trained on 60,000 synthetic states labeled by an explicit rule controller. On 20 matched center-scenario seeds (1000–1019), mean kills were 17.55 for imitation, 17.60 for the teacher, and 0.95 for random control. These runs use two-tic actions and complete steering policies. Their kill counts must not be merged with or directly ranked against the seven-tic firing-only Bayesian study.




Figure 5: All 60 episodes of the original gameplay benchmark. The imitation policy approximates its teacher on this scenario. This is not a generalization or superiority claim.

Rust versus NumPy score kernel. Both implementations receive identical float64 logits with 32 rows, vocabulary width 151,936 and six candidates. They compute candidate probabilities, full-vocabulary candidate mass, and entropy. Inputs include extreme and equal-valued rows. After three warmups, each implementation runs 11 timed batches. Maximum output difference is 1.11×10^{-18} . Median elapsed time per row was 0.550013 ms for Python/NumPy and 0.203807 ms for Rust release, giving a ratio of 2.699.




Figure 6: One local score-kernel microbenchmark: points are batch time divided by 32; red marks show medians and interquartile ranges, not confidence intervals. This excludes model inference, tokenization, Doom, HTTP, startup and I/O.

This measurement used Linux ARM64 in Docker with a two-CPU limit, Python 3.12.14, NumPy 2.5.3 and Rust 1.94.0. Rust ran before Python, and other Docker services remained active. NumPy

already executes native numerical kernels, so the result compares these implementations, not interpreted arithmetic with compiled arithmetic. It is not evidence of a general Rust advantage or end-to-end agent acceleration. Randomized execution order, additional machines and end-to-end profiling remain necessary for stronger claims.

7 Related work and limitations

Calibration for learned dynamics and control is established research (Malik et al., 2019). The Laplace approximation used here is also established (Daxberger et al., 2021); applying it to this logistic head is not a new Bayesian method. RLCR combines correctness and calibration rewards for reasoning language models (Damani et al., 2025). We do not train or reproduce RLCR, and its results cannot be transferred to our firing-policy experiment.

The initial pilot’s linear features omit weapon cooldown and temporal state; the follow-up adds hand-crafted history but not a learned recurrent state. Its conditional-independence likelihood ignores within-episode correlation, which may invalidate the posterior’s uncertainty interpretation. The lower-bound policy may suppress useful firing, but this study does not isolate the mechanism. Both studies use only two scenarios from one game and five training sets. The initial pilot uses ten evaluation seeds per scenario, and the memory follow-up uses thirty; this remains limited evidence for generality. Neither the eight disclosed language development questions nor a single gameplay GIF is an independent language benchmark. No evaluated connectome-inspired architecture, latent transformer, recurrent world model, semantic-entropy method or conformal decision guarantee is established by this report.

8 Continuation decision and reproducibility

Conditional on a future separately justified protocol, a follow-up could ask whether additional observations or recurrent computation improve the *decision*, conditional on their cost. Forecast learning and action selection should remain separately measurable. A proposed controller would allocate extra work only when estimated decision improvement exceeds its cost, rather than using posterior width alone. This is a research hypothesis, not an implemented contribution in these results. The present continuation gate failed, so this work is deferred.

Required controls include the same head without Bayes, fixed extra computation, scalar probability and entropy gates, and matched-data/compute ensembles. Random audit actions need known propensities and support checks. Development and confirmation splits should be fixed before tuning, with a second environment and an English decision task. Literature review must establish a distinct method or theoretical result before any ICLR novelty claim.

The repository retains the v1 and v2 execution snapshots at commits 4007e64 and f59ca5c, frozen protocols, source hashes, episode summaries, posterior weights and losslessly compressed traces. Hashes of raw traces refer to decompressed bytes. V2 completed in 85.18 workstation seconds; this duration is not a portable throughput benchmark. All five fits had gradient norms below 8×10^{-8} and positive covariance eigenvalues, validating numerical mechanics rather than model specification.

The memory run completed 1,980 new episodes in 412.13 workstation seconds. Its preserved analysis records the failed continuation gate. The browser demo uses pinned Qwen3-0.6B weights through WebLLM and runs separately from these numerical controllers; it is not a live Doom port or a quality benchmark. The historical Python package remains `openjev` for compatibility.

The plotting scripts regenerate figures and tables from preserved JSON, separately from experiment execution. Source and artifact hashes accompany the runs. Invalid v1 observations are not pooled into

valid studies. The renamed repository retains historical artifacts and compatibility commands.

References

- Mehul Damani, Isha Puri, Stewart Slocum, Idan Shenfeld, Leshem Choshen, Yoon Kim, and Jacob Andreas. Beyond binary rewards: Training LMs to reason about their uncertainty, 2025. URL <https://arxiv.org/abs/2507.16806>. arXiv:2507.16806, version 2 revised May 2026.
- Erik Daxberger, Agustinus Kristiadi, Alexander Immer, Runa Eschenhagen, Matthias Bauer, and Philipp Hennig. Laplace redux: Effortless bayesian deep learning. In *Advances in Neural Information Processing Systems*, 2021. URL <https://arxiv.org/abs/2106.14806>.
- Michał Kempka, Marek Wydmuch, Grzegorz Runc, Jakub Toczek, and Wojciech Jaśkowski. ViZ-Doom: A doom-based ai research platform for visual reinforcement learning. In *IEEE Conference on Computational Intelligence and Games*, 2016. URL <https://arxiv.org/abs/1605.02097>.
- Ali Malik, Volodymyr Kuleshov, Jiaming Song, Danny Nemer, Harlan Seymour, and Stefano Ermon. Calibrated model-based deep reinforcement learning. In *Proceedings of the 36th International Conference on Machine Learning*, volume 97, pages 4314–4323. PMLR, 2019. URL <https://proceedings.mlr.press/v97/malik19a.html>.