


NL-LFR training and state-estimation budgets



C100/H128; 12 exposed 100/200 mV DEV records. Same saved targets and standardized-difference scoring.

Known future inputs; no future outputs in inference. Old 10,000-iteration fit remains diagnostic.

Four cells share one timed cache adapter. Candidate error is a reference only; no candidate speed comparison.

Finite capped or stalled context solves remain scoreable. No untouched-test, control or novelty claim.