



Query-centered nonlinear features: a bounded learning pilot

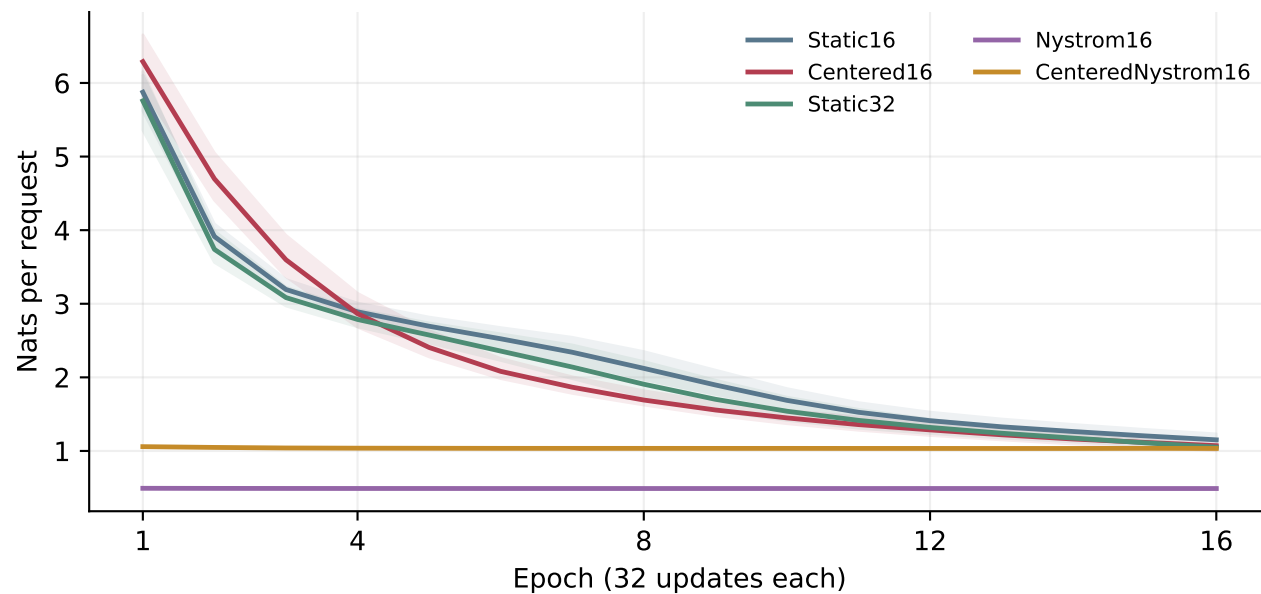
Decision regret under the exact public GP



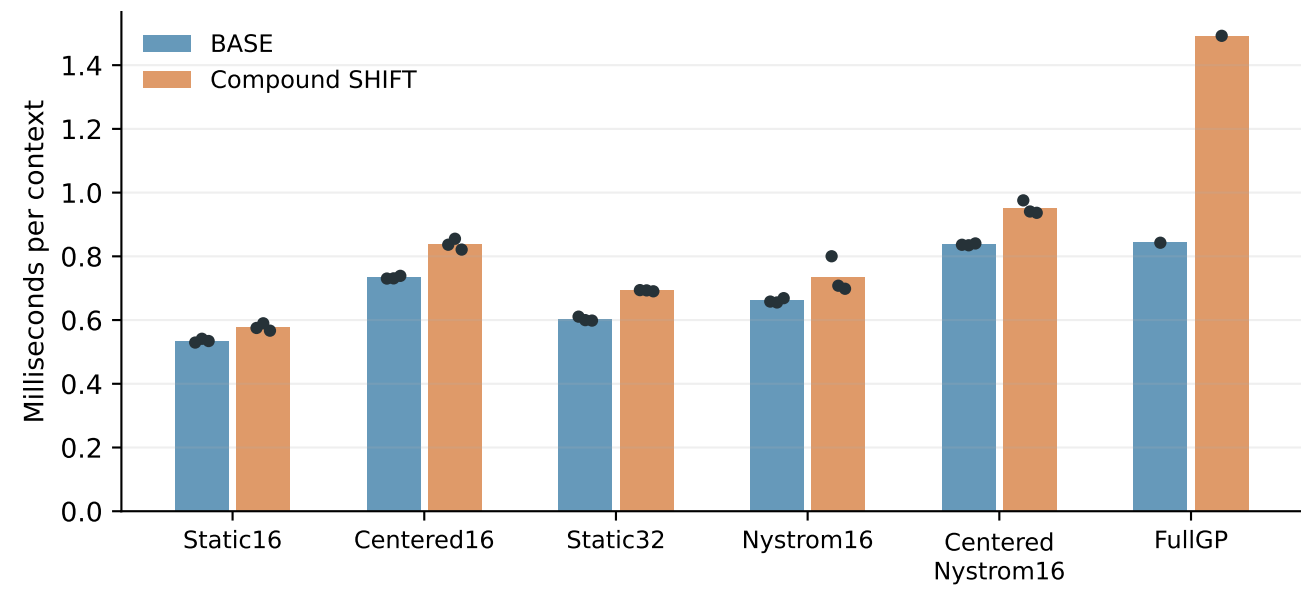
Predictive negative log likelihood



TRAIN NLL: mean and range of three fits



Four-request context latency, including archive work



Metric dots: three cohorts; fitted models average three fits, FullGP uses one reference. Diamonds average cohorts; no ensemble. Static methods cache each archive; centered methods rebuild per request. FullGP uses the known kernel and unoptimized NumPy. Latencies are descriptive measurements on the first context, not an equal-compute comparison or a speed claim.