



Residual covariance with fixed learned kernels

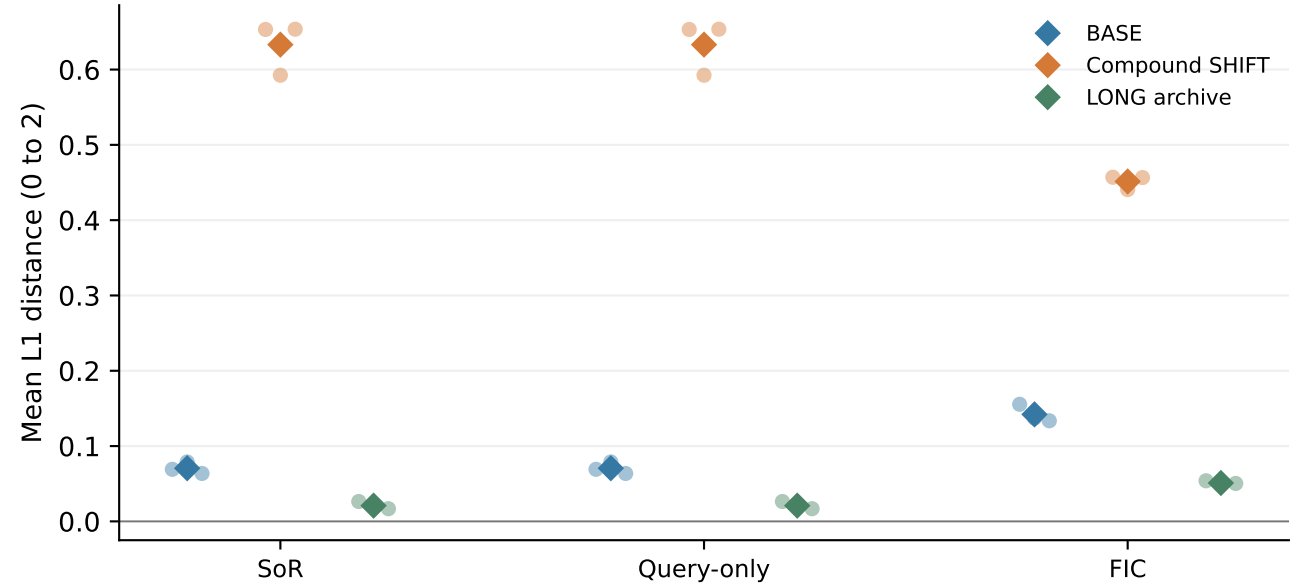
Decision regret under the true public-history GP



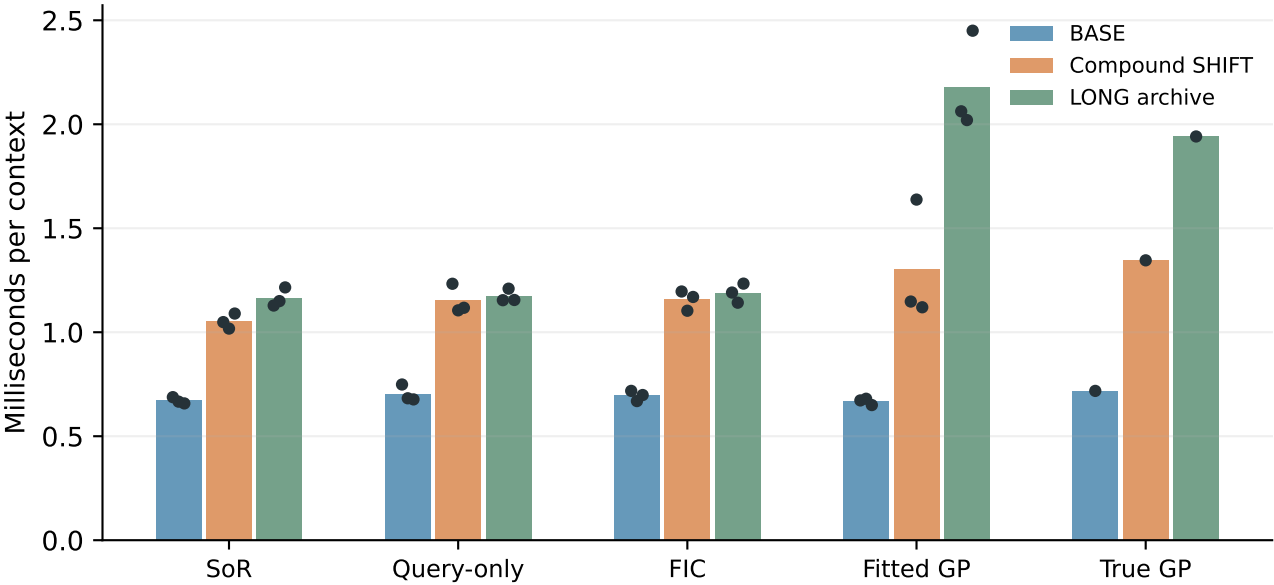
Predictive negative log likelihood



Block-identity weights vs the same fitted full GP



Archive build plus four cached requests



Dots: three independent cohorts; fitted methods average three inherited fits. Diamonds: cohort means, not ensembles.
Fitted GP retains each inherited kernel; True GP knows the generating kernel. No models are trained in this diagnostic.
All modes cache archives. Latency is descriptive on one context per population, not a general speed claim.